

```
# library loading
!pip install pandas
import pandas as pd
```

```
Requirement already satisfied: pandas in c:\users\silom\anaconda3\lib\site-packages (2.2.2)
Requirement already satisfied: numpy>=1.26.0 in c:\users\silom\anaconda3\lib\site-packages (from pandas) (1.26.4)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\silom\anaconda3\lib\site-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\silom\anaconda3\lib\site-packages (from pandas) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in c:\users\silom\anaconda3\lib\site-packages (from pandas) (2023.3)
Requirement already satisfied: six>=1.5 in c:\users\silom\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas) (1.1)
```

```
""" name and code
customer_demographic - cd
transaction_history - th
customer_service - csv
online_activity - oa
churn_status - cs
"""

# uploading sheets
cd = pd.read_excel("Customer_Churn_Data_Large.xlsx", sheet_name = 0)
th = pd.read_excel("Customer_Churn_Data_Large.xlsx", sheet_name = 1)
csv = pd.read_excel("Customer_Churn_Data_Large.xlsx", sheet_name = 2)
oa = pd.read_excel("Customer_Churn_Data_Large.xlsx", sheet_name = 3)
cs = pd.read_excel("Customer_Churn_Data_Large.xlsx", sheet_name = 4)
```

```
#summary and description: customer demographic
print(cd.describe())
print(cd.head())
print("done")
```

	CustomerID	Age
count	1000.00000	1000.00000
mean	500.50000	43.26700
std	288.819436	15.242311
min	1.00000	18.00000
25%	250.75000	30.00000
50%	500.50000	43.00000
75%	750.25000	56.00000
max	1000.00000	69.00000

	CustomerID	Age	Gender	MaritalStatus	IncomeLevel
0	1	62	M	Single	Low
1	2	65	M	Married	Low
2	3	18	M	Single	Low
3	4	21	M	Widowed	Low
4	5	21	M	Divorced	Medium

done

```
#summary and description: transaction history
print(th.describe())
print(th.head())
print("done")
```

	CustomerID	TransactionID	TransactionDate	AmountSpent
count	5054.00000	5054.00000	5054	5054.00000
mean	501.424218	5510.538979	2022-07-01 19:25:37.158686208	250.707351
min	1.00000	1000.00000	2022-01-01 00:00:00	5.180000
25%	251.00000	3242.00000	2022-04-03 00:00:00	127.105000
50%	506.00000	5530.00000	2022-07-01 00:00:00	250.525000
75%	749.00000	7680.75000	2022-09-29 00:00:00	373.412500
max	1000.00000	9997.00000	2022-12-31 00:00:00	499.860000
std	285.172780	2582.088012	NaN	142.250838

	CustomerID	TransactionID	TransactionDate	AmountSpent	ProductCategory
0	1	7194	2022-03-27	416.50	Electronics
1	2	7250	2022-08-08	54.96	Clothing
2	2	9660	2022-07-25	197.50	Electronics
3	2	2998	2022-01-25	101.31	Furniture
4	2	1228	2022-07-24	397.37	Clothing

done

```
#summary and description: customer service
print(csv.describe())
```

```

print(csv.head())
print("done")

      CustomerID  InteractionID          InteractionDate
count  1002.000000    1002.000000           1002
mean   485.209581    5952.887226  2022-07-02 19:28:22.994011904
min    1.000000     2015.000000           2022-01-01 00:00:00
25%   238.250000     3991.500000           2022-04-07 00:00:00
50%   474.500000     5911.500000           2022-07-02 12:00:00
75%   735.750000     7908.250000           2022-09-30 00:00:00
max   995.000000     9997.000000           2022-12-30 00:00:00
std   287.030259    2305.819681             NaN
   CustomerID  InteractionID InteractionDate InteractionType ResolutionStatus
0            1          6363  2022-03-31       Inquiry      Resolved
1            2          3329  2022-03-17       Inquiry      Resolved
2            3          9976  2022-08-24       Inquiry      Resolved
3            4          7354  2022-11-18       Inquiry      Resolved
4            4          5393  2022-07-03       Inquiry  Unresolved
done

```

```

#summary and description: online activity
print(oa.describe())
print(oa.head())
print("done")

```

```

      CustomerID      LastLoginDate  LoginFrequency
count  1000.000000                 1000    1000.000000
mean   500.500000  2023-07-05 21:28:48        25.912000
min    1.000000  2023-01-01 00:00:00        1.000000
25%   250.750000  2023-04-08 00:00:00       13.750000
50%   500.500000  2023-07-10 12:00:00       27.000000
75%   750.250000  2023-10-01 06:00:00       38.000000
max   1000.000000 2023-12-31 00:00:00       49.000000
std   288.819436                  NaN      14.055953
   CustomerID LastLoginDate  LoginFrequency ServiceUsage
0            1  2023-10-21              34  Mobile App
1            2  2023-12-05               5  Website
2            3  2023-11-15               3  Website
3            4  2023-08-25               2  Website
4            5  2023-10-27              41  Website
done

```

```

#summary and description: churn status
print(cs.describe())
print(cs.head())
print("done")

```

```

      CustomerID  ChurnStatus
count  1000.000000  1000.000000
mean   500.500000    0.204000
std   288.819436    0.403171
min    1.000000    0.000000
25%   250.750000    0.000000
50%   500.500000    0.000000
75%   750.250000    0.000000
max   1000.000000   1.000000
   CustomerID ChurnStatus
0            1          0
1            2          1
2            3          0
3            4          0
4            5          0
done

```

```

# python packages for data visualisation
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Set visualization style
sns.set_style("whitegrid")
plt.style.use('ggplot')
%matplotlib inline

# Suppress warnings for cleaner output
import warnings
warnings.filterwarnings('ignore')

```

```
#visualising raw data - intial EDA

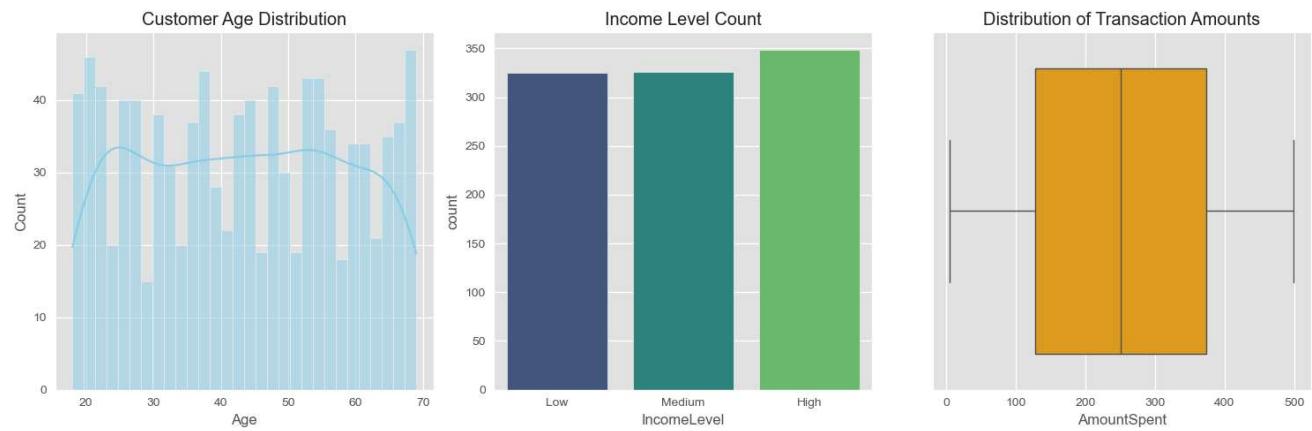
# Set up a figure with multiple subplots
plt.figure(figsize=(15, 5))

# 1. Age Distribution (Are there many young professionals?)
plt.subplot(1, 3, 1)
sns.histplot(cd['Age'], bins=30, kde=True, color='skyblue')
plt.title('Customer Age Distribution')

# 2. Income Levels
plt.subplot(1, 3, 2)
sns.countplot(x='IncomeLevel', data=cd, order=['Low', 'Medium', 'High'], palette='viridis')
plt.title('Income Level Count')

# 3. Transaction Amounts (Checking for outliers)
plt.subplot(1, 3, 3)
sns.boxplot(x=th['AmountSpent'], color='orange')
plt.title('Distribution of Transaction Amounts')

plt.tight_layout()
plt.show()
```



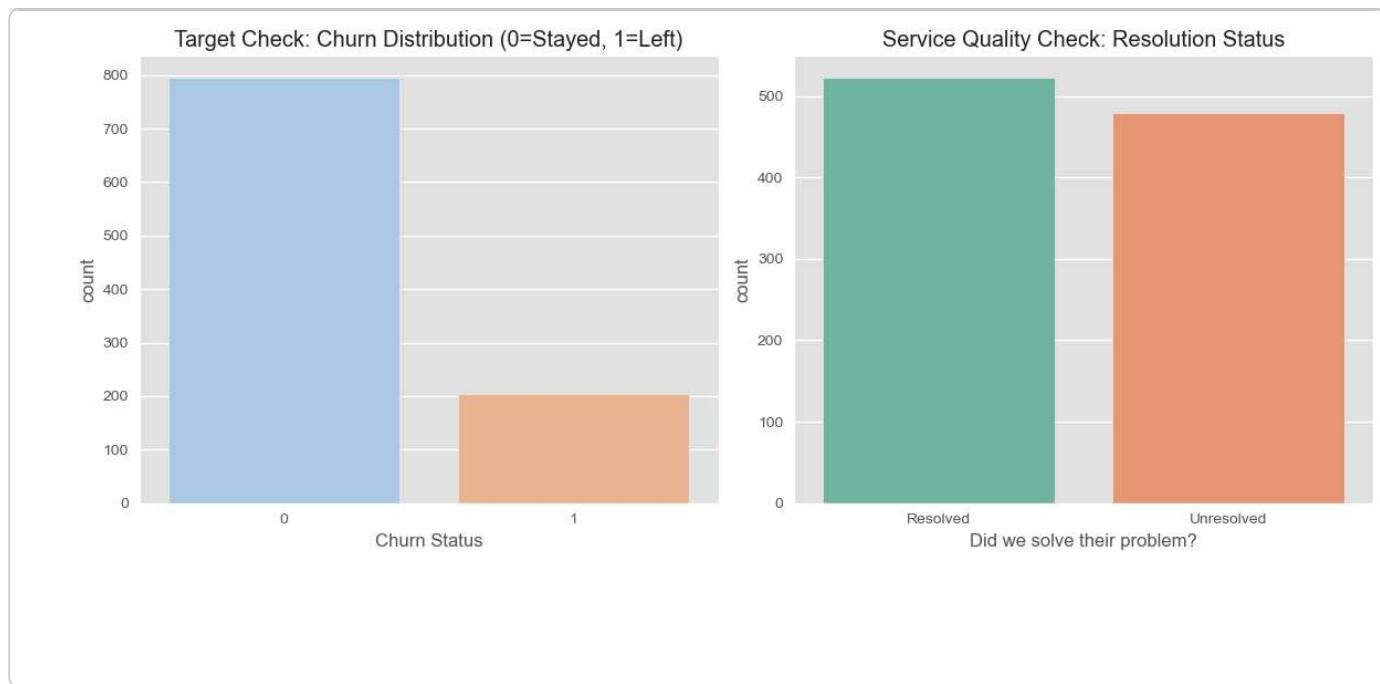
```
#churn and service visualisation

# Create a new figure for the missing categorical checks
plt.figure(figsize=(12, 5))

# 1. Check the Target: Churn Status
plt.subplot(1, 2, 1)
# We use the raw churn_df here
sns.countplot(x='ChurnStatus', data=cs, palette='pastel')
plt.title('Target Check: Churn Distribution (0=Stayed, 1=Left)')
plt.xlabel('Churn Status')

# 2. Check the Pain Point: Service Resolution
plt.subplot(1, 2, 2)
# We use the raw service_df here
sns.countplot(x='ResolutionStatus', data=csv, palette='Set2')
plt.title('Service Quality Check: Resolution Status')
plt.xlabel('Did we solve their problem?')

plt.tight_layout()
plt.show()
```



```
# Feature Engineering - Aggregating Transaction Data
```

```
# Group by CustomerID to get a single summary row per customer
trans_agg = th.groupby('CustomerID').agg(
    Total_Spend=('AmountSpent', 'sum'), # Total money spent
    Avg_Transaction_Amt=('AmountSpent', 'mean'), # Average purchase size
    Transaction_Count=('TransactionID', 'count'), # How active are they?
    Latest_Transaction=('TransactionDate', 'max') # When did they last buy?
).reset_index()

print("Transactions Aggregated:")
display(trans_agg.head())
```

Transactions Aggregated:

CustomerID	Total_Spend	Avg_Transaction_Amt	Transaction_Count	Latest_Transaction
0	1	416.50	416.50000	1
1	2	1547.42	221.06000	7
2	3	1702.98	283.83000	6
3	4	917.29	183.45800	5
4	5	2001.49	250.18625	8

```
# Feature Engineering - Customer Service Interaction Aggregation
```

```
# Group by CustomerID to summarize service history
service_agg = csv.groupby('CustomerID').agg(
    Total_Calls=('InteractionID', 'count'),
    # Calculate how many calls ended as 'Unresolved'
    Unresolved_Calls=('ResolutionStatus', lambda x: (x == 'Unresolved').sum())
).reset_index()

print("Service Interactions Aggregated:")
display(service_agg.head())
```

Service Interactions Aggregated:

CustomerID	Total_Calls	Unresolved_Calls
0	1	1
1	2	1
2	3	1
3	4	2
4	6	1

```
# Master Data Table - merging all aggregated tables on the customer id attribute

# 1. Base Table: Demographics
df = cd.copy()

# 2. Join Churn Status
df = pd.merge(df, cs, on='CustomerID', how='left')

# 3. Join Online Activity
df = pd.merge(df, oa, on='CustomerID', how='left')

# 4. Join Aggregated Transactions
df = pd.merge(df, th, on='CustomerID', how='left')

# 5. Join Aggregated Service Data
df = pd.merge(df, csv, on='CustomerID', how='left')

# Fill NaN values with 0
# (If a customer isn't in the transaction table, it means they spent 0, not that the data is missing)
df.fillna({
    'Total_Spend': 0,
    'Avg_Transaction_Amt': 0,
    'Transaction_Count': 0,
    'Total_Calls': 0,
    'Unresolved_Calls': 0
}, inplace=True)

print(f"Final Master Dataset Shape: {df.shape}")
print(display(df.head()))
print(df.describe())
print(df.shape)
print("done")
```

Final Master Dataset Shape: (6812, 17)

	CustomerID	Age	Gender	MaritalStatus	IncomeLevel	ChurnStatus	LastLoginDate	LoginFrequency	ServiceUsage	TransactionID
0	1	62	M	Single	Low	0	2023-10-21	34	Mobile App	7194
1	2	65	M	Married	Low	1	2023-12-05	5	Website	7250
2	2	65	M	Married	Low	1	2023-12-05	5	Website	9660
3	2	65	M	Married	Low	1	2023-12-05	5	Website	2998
4	2	65	M	Married	Low	1	2023-12-05	5	Website	1228
None										
	CustomerID	Age	ChurnStatus		LastLoginDate	\				
count	6812.000000	6812.000000	6812.000000		6812					
mean	500.169260	43.274516	0.198473	2023-07-06 04:56:47.633587712						
min	1.000000	18.000000	0.000000	2023-01-01 00:00:00						
25%	247.000000	30.000000	0.000000	2023-04-12 00:00:00						
50%	505.500000	44.000000	0.000000	2023-07-13 00:00:00						
75%	750.000000	56.000000	0.000000	2023-10-01 00:00:00						
max	1000.000000	69.000000	1.000000	2023-12-31 00:00:00						
std	286.704642	15.286788	0.398880		Nan					
	LoginFrequency	TransactionID		TransactionDate	\					
count	6812.000000	6812.000000		6812						
mean	25.724310	5497.323253	2022-07-01 02:11:29.136817408							
min	1.000000	1000.000000		2022-01-01 00:00:00						
25%	14.000000	3223.500000		2022-04-02 00:00:00						
50%	26.000000	5515.000000		2022-07-01 00:00:00						
75%	38.000000	7675.250000		2022-09-30 00:00:00						
max	49.000000	9997.000000		2022-12-31 00:00:00						
std	14.062032	2584.768541		Nan						
	AmountSpent	InteractionID		InteractionDate	\					
count	6812.000000	5204.000000		5204						
mean	251.620527	5921.861261	2022-07-01 16:31:10.561106944							
min	5.180000	2015.000000		2022-01-01 00:00:00						
25%	127.100000	3873.250000		2022-04-07 00:00:00						
50%	251.845000	5903.000000		2022-07-02 00:00:00						
75%	375.280000	7851.000000		2022-09-25 00:00:00						
max	499.860000	9997.000000		2022-12-30 00:00:00						
std	142.901693	2332.331260		Nan						

(6812, 17)

done

```
#MERGE AND Feature Engineering - Converting date format to number for machine learning part. models don't read dates but number

# --- PART 1: RE-MERGE (Ensuring columns exist) ---
# 1. Base Table: Demographics
df = cd.copy()

# 2. Join Churn Status
df = pd.merge(df, cs, on='CustomerID', how='left')

# 3. Join Online Activity (Contains 'LastLoginDate')
df = pd.merge(df, oa, on='CustomerID', how='left')

# 4. Join Aggregated Transactions (Contains 'Latest_Transaction')
df = pd.merge(df, th, on='CustomerID', how='left')

# 5. Join Aggregated Service Data
df = pd.merge(df, csv, on='CustomerID', how='left')

# Fill NaN values for numerical columns only
df.fillna({
    'Total_Spend': 0,
    'Avg_Transaction_Amt': 0,
    'Transaction_Count': 0,
    'Total_Calls': 0,
    'Unresolved_Calls': 0
}, inplace=True)

# --- PART 2: DATE ENGINEERING (The fix) ---

# 1. Convert to datetime (using errors='coerce' to handle any messy data safely)
df['LastLoginDate'] = pd.to_datetime(df['LastLoginDate'], errors='coerce')
df['TransactionDate'] = pd.to_datetime(df['TransactionDate'], errors='coerce')
df['InteractionDate'] = pd.to_datetime(df['InteractionDate'], errors='coerce')

# 2. Define Reference Date (Simulating "Today")
# We use the latest date in the entire dataset + 1 day
reference_date = df['LastLoginDate'].max() + pd.Timedelta(days=1)
print(f"Reference Date (Today): {reference_date}")

# 3. Calculate Days Since Last Event
df['Days_Since_Login'] = (reference_date - df['LastLoginDate']).dt.days
df['Days_Since_Transaction'] = (reference_date - df['TransactionDate']).dt.days
df['Days_Since_Interaction'] = (reference_date - df['InteractionDate']).dt.days

# 4. Handle NaNs for dates
# If Days_Since_Login is NaN, it means they never logged in. We set it to a high number (e.g., 365).
df['Days_Since_Login'].fillna(365, inplace=True)
df['Days_Since_Transaction'].fillna(365, inplace=True)
df['Days_Since_Interaction'].fillna(365, inplace=True)

# 5. Drop the original date columns (Cleanup)
df.drop(columns=['LastLoginDate', 'TransactionDate', 'InteractionDate'], inplace=True)

print("--- Date Engineering Successful ---")
print("New features 'Days_Since_Login' and 'Days_Since_Transaction', 'Days_Since_Interaction' created.")
display(df[['CustomerID', 'Days_Since_Login', 'Days_Since_Transaction', 'Days_Since_Interaction']].head())

```

Reference Date (Today): 2024-01-01 00:00:00  
--- Date Engineering Successful ---  
New features 'Days\_Since\_Login' and 'Days\_Since\_Transaction', 'Days\_Since\_Interaction' created.

CustomerID	Days_Since_Login	Days_Since_Transaction	Days_Since_Interaction
0	1	72	645
1	2	27	511
2	2	27	525
3	2	27	706
4	2	27	526

```
# Correlation with churn check

# Select only numeric columns for correlation
numeric_df = df.select_dtypes(include=[np.number])

# Calculate correlation with ChurnStatus
```

```

correlation = numeric_df.corr()['ChurnStatus'].sort_values(ascending=False)

print("Correlation with Churn Status:")
print(correlation)

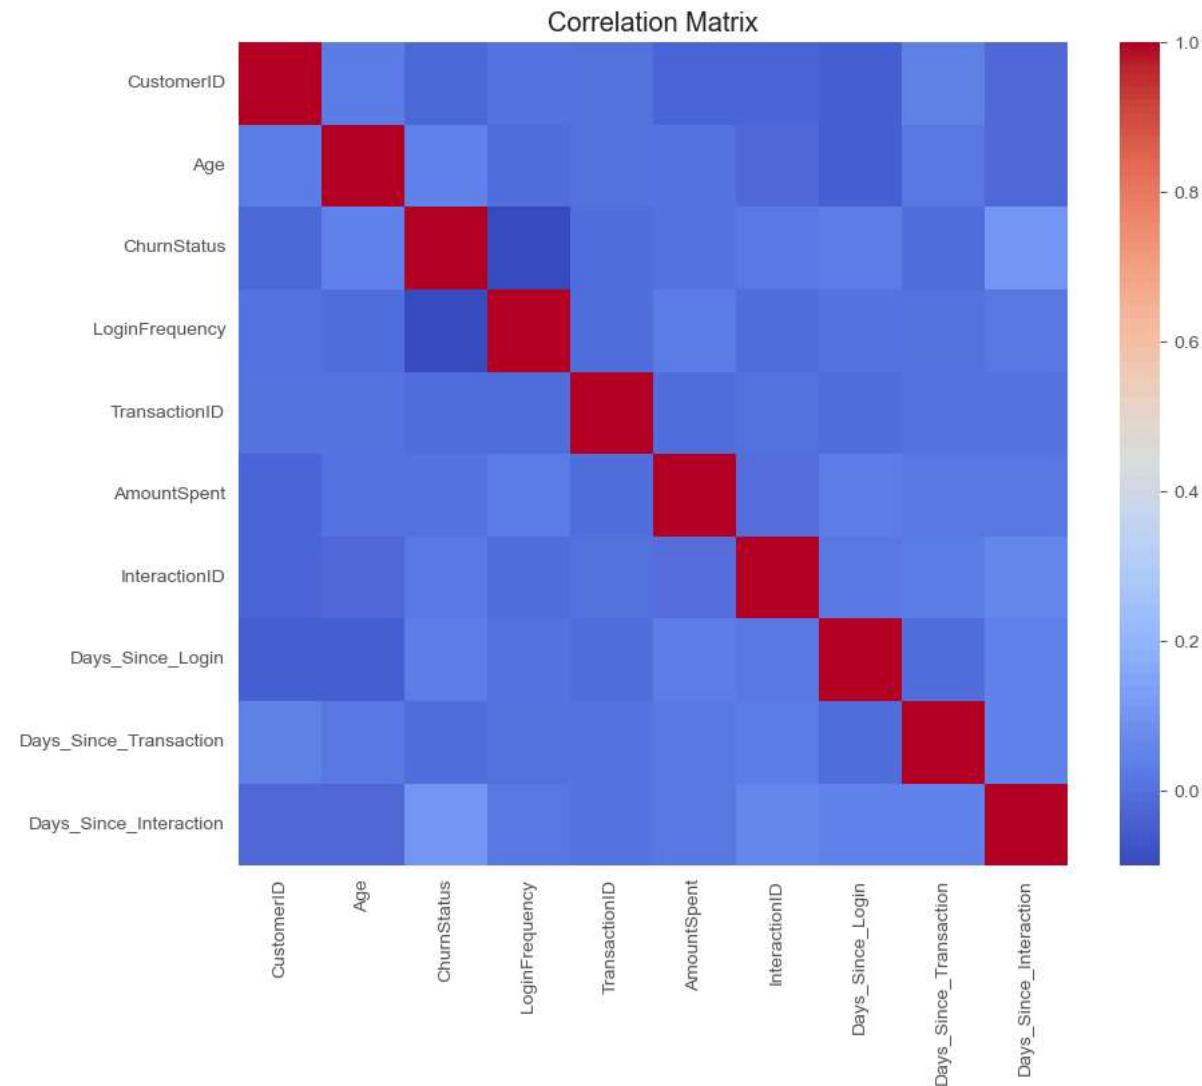
# Visual Heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(numeric_df.corr(), cmap='coolwarm', annot=False)
plt.title("Correlation Matrix")
plt.show()

```

```

Correlation with Churn Status:
ChurnStatus      1.000000
Days_Since_Interaction  0.093400
Age              0.045048
Days_Since_Login   0.031073
InteractionID    0.023130
AmountSpent       0.005113
Days_Since_Transaction -0.006527
TransactionID    -0.008736
CustomerID        -0.025158
LoginFrequency    -0.100391
Name: ChurnStatus, dtype: float64

```



```

# Final Processing for machine learning and further parts in follow up steps

# 1. One-Hot Encode Categorical Variables
# This converts 'Gender' -> 'Gender_Male' (0 or 1)
df_model = pd.get_dummies(df, columns=['Gender', 'MaritalStatus', 'IncomeLevel', 'ServiceUsage'], drop_first=True)

# 2. Drop ID column (not useful for prediction)
df_model.drop(columns=['CustomerID'], inplace=True)

print("Final Data Ready for Modelling:")

```

```
display(df_model.head())

# 3. Save to CSV
df_model.to_csv('cleaned_lloyds_data.csv', index=False)
print("File 'cleaned_lloyds_data.csv' saved successfully!")
```

Final Data Ready for Modelling:

	Age	ChurnStatus	LoginFrequency	TransactionID	AmountSpent	ProductCategory	InteractionID	InteractionType	ResolutionStatus
0	62	0	34	7194	416.50	Electronics	6363.0	Inquiry	Resolved
1	65	1	5	7250	54.96	Clothing	3329.0	Inquiry	Resolved
2	65	1	5	9660	197.50	Electronics	3329.0	Inquiry	Resolved
3	65	1	5	2998	101.31	Furniture	3329.0	Inquiry	Resolved
4	65	1	5	1228	397.37	Clothing	3329.0	Inquiry	Resolved

File 'cleaned\_lloyds\_data.csv' saved successfully!

```
#Visualisation - final

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# =====
# 1. SETUP AND CONFIGURATION
# =====

TARGET_COL = 'ChurnStatus'

# List of remaining numerical features (Confirmed in previous interaction)
INTENDED_NUMERICAL_COLS = [
    'Age',
    'LoginFrequency',
    'AmountSpent',
    'Days_Since_Login',
    'Days_Since_Transaction',
    'Days_Since_Interaction'
]

# The CORRECT list of one-hot encoded binary columns
INTENDED_BINARY_COLS = [
    'Gender_M',
    'MaritalStatus_Married',
    'MaritalStatus_Single',
    'MaritalStatus_Widowed',
    'IncomeLevel_Low',
    'IncomeLevel_Medium',
    'ServiceUsage_OnlineBanking', # <-- CORRECTED
    'ServiceUsage_Website'      # <-- CORRECTED
]
# =====

# Load the processed DataFrame
try:
    df_model = pd.read_csv('cleaned_lloyds_data.csv')
except FileNotFoundError:
    print("Error: 'cleaned_lloyds_data.csv' not found. Please check your file path and ensure the file exists.")
    exit()

# --- ROBUSTNESS CHECK: Filter columns to only include those present in the DataFrame ---
df_cols = df_model.columns.tolist()
NUMERICAL_COLS = [col for col in INTENDED_NUMERICAL_COLS if col in df_cols]
BINARY_COLS = [col for col in INTENDED_BINARY_COLS if col in df_cols]

if TARGET_COL not in df_cols:
    print(f"\nFATAL ERROR: Target column '{TARGET_COL}' is missing from the data.")
    exit()

print(f"Loaded DataFrame with {len(df_cols)} columns.")
print(f"Numerical columns used: {NUMERICAL_COLS}")
print(f"Binary columns used: {BINARY_COLS}")
# =====
```

```

# =====
# 2. Target Variable Class Balance (Bar Plot)
# =====

plt.figure(figsize=(6, 4))
sns.countplot(x=TARGET_COL, data=df_model, palette='viridis')
plt.title(f'Distribution of the Target Variable ({TARGET_COL})')
plt.xlabel(TARGET_COL)
plt.ylabel('Count')
plt.xticks([0, 1], ['No Churn (0)', 'Churn (1)'])
plt.savefig('2_target_balance_bar_plot.png')
plt.show()
plt.close()

# =====
# 3. Correlation Heatmap
# =====
# Identify multicollinearity and feature correlation with ChurnStatus.

relevant_cols = [TARGET_COL] + NUMERICAL_COLS + BINARY_COLS
corr_matrix = df_model[relevant_cols].corr()

plt.figure(figsize=(14, 12))
sns.heatmap(corr_matrix,
            annot=True,
            fmt=".2f",
            cmap='coolwarm',
            cbar=True,
            linewidths=.5,
            linecolor='black')
plt.title(f'Correlation Heatmap of Features vs. {TARGET_COL}')
plt.savefig('3_correlation_heatmap.png')
plt.show()
plt.close()

# =====
# 4. Numerical Feature Distributions (Histograms)
# =====
# Check for skewness and general distribution shape of continuous features.

n_plots = len(NUMERICAL_COLS)
n_cols = 3
n_rows = (n_plots + n_cols - 1) // n_cols

plt.figure(figsize=(15, 5 * n_rows))
for i, col in enumerate(NUMERICAL_COLS):
    plt.subplot(n_rows, n_cols, i + 1)
    sns.histplot(df_model[col], kde=True, bins=30, edgecolor='black')
    plt.title(f'Distribution of {col}')
    plt.xlabel(col)
    plt.ylabel('Frequency')

plt.tight_layout()
plt.savefig('4_numerical_distributions.png')
plt.show()
plt.close()

# =====
# 5. Binary Features vs. Target (Exit Rate Bar Plots)
# =====
# Visualize how the one-hot encoded groups influence the churn rate.

n_plots = len(BINARY_COLS)
n_cols = 3
n_rows = (n_plots + n_cols - 1) // n_cols

plt.figure(figsize=(15, 5 * n_rows))
for i, col in enumerate(BINARY_COLS):
    # Calculate the mean of the target column (ChurnStatus) for each binary feature class (0 or 1)
    temp_df = df_model.groupby(col)[TARGET_COL].mean().reset_index()

    plt.subplot(n_rows, n_cols, i + 1)

```

```
sns.barplot(x=col, y=TARGET_COL, data=temp_df, palette='pastel')

plt.title(f'Churn Rate by {col}')
plt.xlabel(col)
plt.ylabel(f'Mean {TARGET_COL} (Churn Rate)')
plt.xticks([0, 1], ['No', 'Yes'])

plt.tight_layout()
plt.savefig('5_binary_vs_target_proportions.png')
plt.show()
plt.close()
```

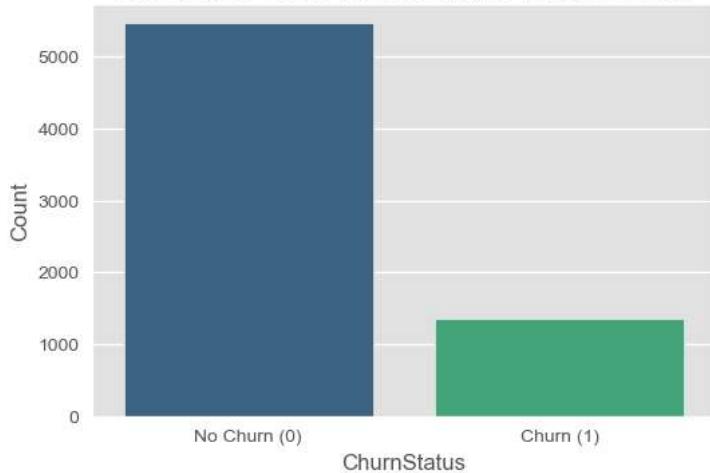


Loaded DataFrame with 20 columns.

Numerical columns used: ['Age', 'LoginFrequency', 'AmountSpent', 'Days\_Since\_Login', 'Days\_Since\_Transaction', 'Days\_Since\_Interaction', 'Gender\_M', 'MaritalStatus\_Married', 'MaritalStatus\_Single', 'MaritalStatus\_Widowed', 'IncomeLevel\_Low', 'IncomeLevel\_Medium', 'IncomeLevel\_High', 'ServiceUsage\_Website', 'ServiceUsage\_Mobile', 'ServiceUsage\_Phone', 'ServiceUsage\_Email', 'ServiceUsage\_Facebook', 'ServiceUsage\_Gmail', 'ServiceUsage\_YouTube', 'ServiceUsage\_Twitter', 'ServiceUsage\_Snapchat', 'ServiceUsage\_Instagram']

Binary columns used: ['Gender\_M', 'MaritalStatus\_Married', 'MaritalStatus\_Single', 'MaritalStatus\_Widowed', 'IncomeLevel\_Low', 'IncomeLevel\_Medium', 'IncomeLevel\_High', 'ServiceUsage\_Website', 'ServiceUsage\_Mobile', 'ServiceUsage\_Phone', 'ServiceUsage\_Email', 'ServiceUsage\_Facebook', 'ServiceUsage\_Gmail', 'ServiceUsage\_YouTube', 'ServiceUsage\_Twitter', 'ServiceUsage\_Snapchat', 'ServiceUsage\_Instagram']

### Distribution of the Target Variable (ChurnStatus)



### Correlation Heatmap of Features vs. ChurnStatus

